

# Project : Churn rate prediction

Company target

- Churn rate is 1.7%
- retention rate is 80%

## 1. Importing frameworks and Data preparation

According to current situation, how to do proactive action to reduce churn rate, especially, in postpaid loyalty churning.

- Filter postpaid customers.

Loyalty customer is customers those who stayed with True more than 3 years.

- Filter loyalty customers.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv('customer_dataset.csv')
df_postpaid_loyalty = df[(df['ACCT_TYPE']=='Postpaid') & (df['AOU_DAY']>1095)]
df_postpaid_loyalty.head()
```

```
Out[1]:
```

	MNTH_ID	SUBR_ID	ACCT_TYPE	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_BRAND	DVC_MODEL	DVC_GRP	DVC_CLASS	...	VC_DO
0	202009	1879	Postpaid	1099.0	10519	467	SAMSUNG	GALAXY A10 [SM-A105GDS]	Smartphone	Phone	...	
1	202010	1879	Postpaid	1099.0	10550	498	SAMSUNG	GALAXY A10 [SM-A105GDS]	Smartphone	Phone	...	
2	202011	1879	Postpaid	1099.0	10580	528	SAMSUNG	GALAXY A10 [SM-A105GDS]	Smartphone	Phone	...	

	MNTH_ID	SUBR_ID	ACCT_TYPE	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_BRAND	DVC_MODEL	DVC_GRP	DVC_CLASS	...	VC_DO
3	202012	1879	Postpaid	1099.0	10611	559	SAMSUNG	GALAXY A10 [SM-A105GDS]	Smartphone	Phone	...	
4	202009	7739	Postpaid	999.0	10403	110	SAMSUNG	GALAXY A6 [SM-A600GDS]	Smartphone	Phone	...	

5 rows × 30 columns

In [2]:

```
df_postpaid_loyalty.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 180133 entries, 0 to 211339
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   MNTH_ID                               180133 non-null int64
1   SUBR_ID                               180133 non-null int64
2   ACCT_TYPE                             180133 non-null object
3   CURR_MAIN_PKG_FEE                     180133 non-null float64
4   AOU_DAY                               180133 non-null int64
5   AOU_DVC                               180133 non-null object
6   DVC_BRAND                             179994 non-null object
7   DVC_MODEL                             179994 non-null object
8   DVC_GRP                               180133 non-null object
9   DVC_CLASS                             180133 non-null object
10  DVC_SUPPORT                           178599 non-null object
11  MOST_USED_4W_REGION                   180133 non-null object
12  MOST_USED_4W_PRVNC                   180133 non-null object
13  MOST_USED_4W_AMPHUR                   180133 non-null object
14  MOST_USED_4W_TUMBON                   180133 non-null object
15  DTAC_RWRD_SEGMENT                     180133 non-null object
16  Churn_Flag                            180133 non-null object
17  REVENUE                               180133 non-null float64
18  AVG3M_REVENUE                         180133 non-null float64
19  VC_DOM_MOU                            180133 non-null int64
20  VC_DOM_CNT                            180133 non-null int64
21  AVG3M_VC_DOM_MOU                      180133 non-null float64
22  AVG3M_VC_DOM_CNT                      180133 non-null float64
23  DATA_MB                              180133 non-null float64
24  AVG3M_DATA_MB                        180133 non-null float64
25  CIN_CALLCNT                           180133 non-null int64
```

```

26 VOC_ACTIVATEDAY      180133 non-null  int64
27 DATA_ACTIVATEDAY     180133 non-null  int64
28 PM_PMMTHDCOMMON       180133 non-null  object
29 PCT_CALL_DROP         180133 non-null  float64
dtypes: float64(8), int64(8), object(14)
memory usage: 42.6+ MB

```

```
In [3]: df_postpaid_loyalty.describe()
```

```

Out[3]:
```

	MNTH_ID	SUBR_ID	CURR_MAIN_PKG_FEE	AOU_DAY	REVENUE	AVG3M_REVENUE	VC_DOM_MOU	VC_DOM_CNT	AVG3M_REVENUE
<b>count</b>	180133.000000	1.801330e+05	180133.000000	180133.000000	180133.000000	180133.000000	180133.000000	180133.000000	180133.000000
<b>mean</b>	202010.500230	1.239937e+08	522.164028	3672.153181	532.613637	541.229503	188.340448	78.756186	532.613637
<b>std</b>	1.116564	8.887744e+07	312.603697	1920.463407	345.475348	326.534668	314.279055	115.979303	345.475348
<b>min</b>	202009.000000	1.879000e+03	0.000000	1096.000000	-1520.064500	-406.688200	0.000000	0.000000	-1520.064500
<b>25%</b>	202010.000000	3.039002e+07	299.000000	2021.000000	299.000000	309.333300	41.000000	20.000000	299.000000
<b>50%</b>	202010.000000	1.195342e+08	449.000000	3253.000000	463.500000	482.333300	107.000000	48.000000	463.500000
<b>75%</b>	202011.000000	2.146236e+08	699.000000	5118.000000	699.000000	699.000000	225.000000	96.000000	699.000000
<b>max</b>	202012.000000	2.694149e+08	5607.480000	10700.000000	17093.450000	7339.618700	19116.000000	6185.000000	17093.450000

- Drop features, which are not useful.

```
In [4]: df_fs = df_postpaid_loyalty.drop([
        'MNTH_ID', 'SUBR_ID', 'ACCT_TYPE', 'DVC_BRAND', 'DVC_MODEL', 'MOST_USED_4W_PRVNC', 'MOST_USED_4W_AMPHUR', 'MOST_USED_4W_REGION',
        ], axis=1)
```

```
In [5]: df_fs.head(5)
```

```

Out[5]:
```

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_GRP	DVC_CLASS	DVC_SUPPORT	MOST_USED_4W_REGION	DTAC_RWRD_SEGMENT	Churn
<b>0</b>	1099.0	10519	467	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Cust
<b>1</b>	1099.0	10550	498	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Cust

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_GRP	DVC_CLASS	DVC_SUPPORT	MOST_USED_4W_REGION	DTAC_RWRD_SEGMENT	Churn
2	1099.0	10580	528	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Cust
3	1099.0	10611	559	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Cust
4	999.0	10403	110	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Cust

5 rows × 22 columns

In [6]:

```
df_fs.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 180133 entries, 0 to 211339
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CURR_MAIN_PKG_FEE     180133 non-null float64
1   AOU_DAY               180133 non-null int64
2   AOU_DVC              180133 non-null object
3   DVC_GRP              180133 non-null object
4   DVC_CLASS            180133 non-null object
5   DVC_SUPPORT          178599 non-null object
6   MOST_USED_4W_REGION  180133 non-null object
7   DTAC_RWRD_SEGMENT    180133 non-null object
8   Churn_Flag           180133 non-null object
9   REVENUE              180133 non-null float64
10  AVG3M_REVENUE        180133 non-null float64
11  VC_DOM_MOU           180133 non-null int64
12  VC_DOM_CNT           180133 non-null int64
13  AVG3M_VC_DOM_MOU     180133 non-null float64
14  AVG3M_VC_DOM_CNT     180133 non-null float64
15  DATA_MB             180133 non-null float64
16  AVG3M_DATA_MB        180133 non-null float64
17  CIN_CALLCNT          180133 non-null int64
18  VOC_ACTIVATEDAY      180133 non-null int64
19  DATA_ACTIVATEDAY    180133 non-null int64
20  PM_PMMTHDCOMMON      180133 non-null object
21  PCT_CALL_DROP        180133 non-null float64
dtypes: float64(8), int64(6), object(8)
memory usage: 31.6+ MB
```

In [7]: `df_fs.columns`

Out[7]: Index(['CURR\_MAIN\_PKG\_FEE', 'AOU\_DAY', 'AOU\_DVC', 'DVC\_GRP', 'DVC\_CLASS',  
'DVC\_SUPPORT', 'MOST\_USED\_4W\_REGION', 'DTAC\_RWRD\_SEGMENT', 'Churn\_Flag',  
'REVENUE', 'AVG3M\_REVENUE', 'VC\_DOM\_MOU', 'VC\_DOM\_CNT',  
'AVG3M\_VC\_DOM\_MOU', 'AVG3M\_VC\_DOM\_CNT', 'DATA\_MB', 'AVG3M\_DATA\_MB',  
'CIN\_CALLCNT', 'VOC\_ACTIVATEDAY', 'DATA\_ACTIVATEDAY', 'PM\_PMMTHDCOMMON',  
'PCT\_CALL\_DROP'],  
dtype='object')

In [8]: `df_fs['AOU_DVC'] = pd.to_numeric(df_fs['AOU_DVC'], errors='coerce')`

- Data cleaning

In [9]: `df_fs.drop_duplicates(inplace=True)  
df_fs.dropna(inplace=True)  
df_fs = df_fs[~df_fs.isin(['No Information']).any(axis=1)]  
df_fs`

Out[9]:

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_GRP	DVC_CLASS	DVC_SUPPORT	MOST_USED_4W_REGION	DTAC_RWRD_SEGMENT
<b>0</b>	1099.0	10519	467.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue
<b>1</b>	1099.0	10550	498.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue
<b>2</b>	1099.0	10580	528.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue
<b>3</b>	1099.0	10611	559.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue
<b>4</b>	999.0	10403	110.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue
...	...	...	...	...	...	...	...	...
<b>211323</b>	499.0	1096	751.0	Smartphone	Phone	4G 2300MHz	BMA	Silver
<b>211327</b>	299.0	1096	22.0	Smartphone	Phone	4G 2300MHz	Northeast	Gold

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_GRP	DVC_CLASS	DVC_SUPPORT	MOST_USED_4W_REGION	DTAC_RWRD_SEGMENT
211331	399.0	1096	762.0	Smartphone	Phone	4G 2300MHz	North	Silver
211335	399.0	1096	333.0	Smartphone	Phone	4G 2300MHz	BMA	Silver
211339	699.0	1096	1095.0	Tablet	Tablet	4G 2300MHz	Central&West	Silver

178503 rows × 22 columns

```
In [10]: numerical_features_list = ['CURR_MAIN_PKG_FEE', 'AOU_DAY', 'AOU_DVC', 'REVENUE', 'AVG3M_REVENUE', 'VC_DOM_MOU', 'VC_DOM_C',
                                'AVG3M_VC_DOM_MOU', 'AVG3M_VC_DOM_CNT', 'DATA_MB', 'AVG3M_DATA_MB', 'CIN_CALLCNT', 'VOC_ACTIVE',
                                'DATA_ACTIVEDAY', 'PCT_CALL_DROP']

for col in df_fs.columns:
    if col not in numerical_features_list:
        print(col, df_fs[col].unique())
    print("-"*50)
```

```
-----
-----
-----
DVC_GRP ['Smartphone' 'Feature Phone' 'Tablet' 'Standalone connectivity device']
-----
DVC_CLASS ['Phone' 'Tablet' 'Connectivity Devices']
-----
DVC_SUPPORT ['4G 2300MHz' '3G 2100MHz' '4G 2100MHz' '2G 1800MHz' '4G 1800MHz'
             '3G 850MHz']
-----
MOST_USED_4W_REGION ['BMA' 'East' 'South' 'Central&West' 'unknown' 'Northeast' 'North']
-----
DTAC_RWRD_SEGMENT ['Platinum Blue' 'Silver' 'Gold' 'Welcome']
-----
Churn_Flag ['01) Normal Customers' '02) Churn Customers']
-----
-----
-----
-----
-----
-----
-----
-----
-----
-----
-----
```

```
-----
-----
-----
-----
PM_PMMTHDCOMMON ['VISA' 'Bank transfer' 'Cash' 'Recurring RC' 'Cheque Counter'
'Recurring DD' 'MASTER' 'Cheque Mail' 'AMEX' 'Debit Card'
'Payment After Adjust' 'Pre to Post' 'JCB']
-----
-----
```

```
In [11]: print(df_fs.isnull().sum())
```

```
CURR_MAIN_PKG_FEE      0
AOU_DAY                0
AOU_DVC                0
DVC_GRP                0
DVC_CLASS              0
DVC_SUPPORT            0
MOST_USED_4W_REGION    0
DTAC_RWRD_SEGMENT      0
Churn_Flag             0
REVENUE                0
AVG3M_REVENUE          0
VC_DOM_MOU             0
VC_DOM_CNT             0
AVG3M_VC_DOM_MOU       0
AVG3M_VC_DOM_CNT       0
DATA_MB                0
AVG3M_DATA_MB          0
CIN_CALLCNT            0
VOC_ACTIVATEDAY        0
DATA_ACTIVATEDAY       0
PM_PMMTHDCOMMON        0
PCT_CALL_DROP          0
dtype: int64
```

```
In [12]: df_fs['Churn_Flag'].value_counts()
```

```
Out[12]: 01) Normal Customers    175494
02) Churn Customers           3009
Name: Churn_Flag, dtype: int64
```

## 2. Exploratory data analysis (EDA)

```
In [13]: df_fs.shape
```

Out[13]: (178503, 22)

In [14]: `df_fs.columns`

Out[14]: Index(['CURR\_MAIN\_PKG\_FEE', 'AOU\_DAY', 'AOU\_DVC', 'DVC\_GRP', 'DVC\_CLASS',  
'DVC\_SUPPORT', 'MOST\_USED\_4W\_REGION', 'DTAC\_RWRD\_SEGMENT', 'Churn\_Flag',  
'REVENUE', 'AVG3M\_REVENUE', 'VC\_DOM\_MOU', 'VC\_DOM\_CNT',  
'AVG3M\_VC\_DOM\_MOU', 'AVG3M\_VC\_DOM\_CNT', 'DATA\_MB', 'AVG3M\_DATA\_MB',  
'CIN\_CALLCNT', 'VOC\_ACTIVATEDAY', 'DATA\_ACTIVATEDAY', 'PM\_PMMTHDCOMMON',  
'PCT\_CALL\_DROP'],  
dtype='object')

In [15]: `df_fs.head()`

Out[15]:

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_GRP	DVC_CLASS	DVC_SUPPORT	MOST_USED_4W_REGION	DTAC_RWRD_SEGMENT	Churn
0	1099.0	10519	467.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Custc
1	1099.0	10550	498.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Custc
2	1099.0	10580	528.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Custc
3	1099.0	10611	559.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Custc
4	999.0	10403	110.0	Smartphone	Phone	4G 2300MHz	BMA	Platinum Blue	01) N Custc

5 rows × 22 columns



In [16]: `df_fs.describe()`

Out[16]:

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	REVENUE	AVG3M_REVENUE	VC_DOM_MOU	VC_DOM_CNT	AVG3M_VC_DOM_M
count	178503.000000	178503.000000	178503.000000	178503.000000	178503.000000	178503.000000	178503.000000	178503.000000
mean	523.807431	3672.628309	627.353025	534.316257	542.878558	188.492485	78.693691	190.697



	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	REVENUE	AVG3M_REVENUE	VC_DOM_MOU	VC_DOM_CNT	AVG3M_VC_DOM_M
<b>std</b>	312.610852	1920.474449	502.803347	344.979079	326.018580	314.517588	115.152689	303.103
<b>min</b>	0.000000	1096.000000	1.000000	-1520.064500	-406.688200	0.000000	0.000000	0.000
<b>25%</b>	299.000000	2021.000000	222.000000	299.000000	313.808600	41.000000	20.000000	48.666
<b>50%</b>	449.000000	3253.000000	521.000000	468.000000	487.661900	108.000000	48.000000	114.000
<b>75%</b>	699.000000	5118.000000	929.000000	699.000000	699.000000	225.000000	96.000000	225.666
<b>max</b>	5607.480000	10700.000000	3684.000000	17093.450000	7339.618700	19116.000000	6185.000000	15133.333

- Understand the distribution of features

In [17]:

```
def plot_histogram(df_fs, column_name):

    plt.figure(figsize=(8, 4))
    sns.histplot(df_fs[column_name], kde=True)
    plt.title(f"Distribution of {column_name}")

    # calculate the mean and median values for the columns
    col_mean = df_fs[column_name].mean()
    col_median = df_fs[column_name].median()


    # add vertical lines for mean and median
    plt.axvline(col_mean, color="red", linestyle="--", label="Mean")
    plt.axvline(col_median, color="green", linestyle="-", label="Median")

    plt.legend()

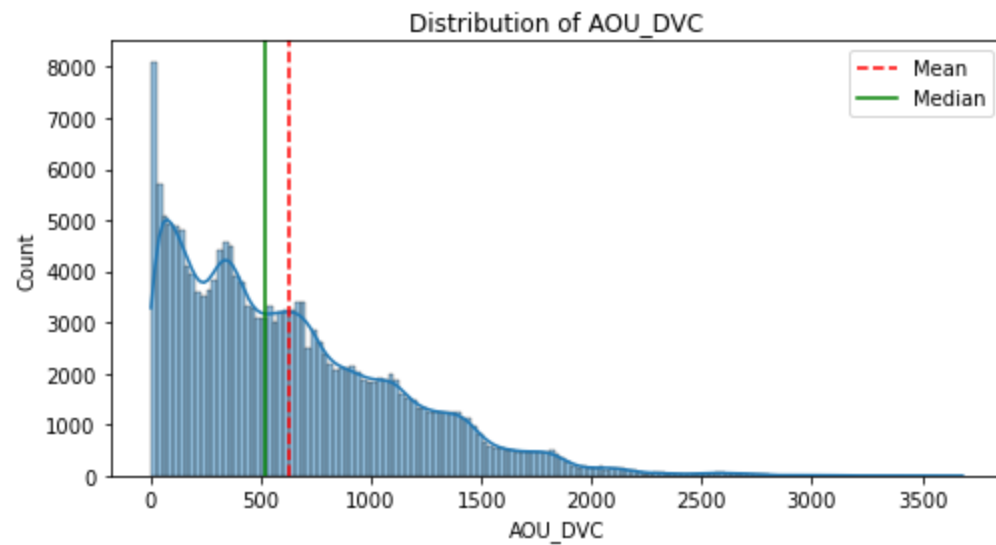
    plt.show()
```

In [18]:


```
plot_histogram(df_fs, 'AOU_DAY')
```




```
In [19]: plot_histogram(df_fs, 'AOU_DVC')
```



```
In [20]: plot_histogram(df_fs, 'CURR_MAIN_PKG_FEE')
```




```
In [21]: plot_histogram(df_fs, 'REVENUE')
```




```
In [22]: df_fs['DTAC_RWRD_SEGMENT'].value_counts().plot(kind='bar')
```

```
Out[22]: <AxesSubplot:>
```




```
In [73]: def plot_boxplot(df_fs, column_name):  
  
    plt.figure(figsize=(5, 8))  
    sns.boxplot(y=df_fs[column_name])  
    plt.title(f"Box Plot of {column_name}")  
    plt.ylabel(column_name)  
    plt.show
```


```
In [74]: plot_boxplot(df_fs, "AOU_DAY")
```




```
In [75]: plot_boxplot(df_fs, "AOU_DVC")
```



```
In [76]: plot_boxplot(df_fs, "CURR_MAIN_PKG_FEE")
```




```
In [77]: plot_boxplot(df_fs, "REVENUE")
```



In [28]:


```
plt.figure(figsize=(12, 6))
sns.heatmap(df_fs[numerical_features_list].corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```





```
In [49]: df_pie = df_fs['Churn_Flag'].value_counts()

plt.figure(figsize=(5, 5))
df_pie.plot(kind='pie', autopct='%1.1f%%', startangle=90, colors=['#ff9999', '#99ff99'], fontsize = 12)
plt.title('', size = 20)
plt.show()
```



### 3. Model conducting

```
In [29]: df_fs['Churn_Flag'] = df_fs['Churn_Flag'].replace({'02) Churn Customers': 1, '01) Normal Customers': 0})
df_fs['Churn_Flag'].value_counts()
```

```
Out[29]: 0    175494
         1     3009
         Name: Churn_Flag, dtype: int64
```

- Label encoding for categorical features

```
In [30]: object_columns = df_fs.select_dtypes(include="object").columns
object_columns
```

```
Out[30]: Index(['DVC_GRP', 'DVC_CLASS', 'DVC_SUPPORT', 'MOST_USED_4W_REGION',
               'DTAC_RWRD_SEGMENT', 'PM_PMMTHDCOMMON'],
              dtype='object')
```

```
In [31]: from sklearn.preprocessing import LabelEncoder
import pickle

encoders = {}
```

```

for column in object_columns:
    label_encoder = LabelEncoder()
    df_fs[column] = label_encoder.fit_transform(df_fs[column])
    encoders[column] = label_encoder

# save the encoders to a pickle file
with open("encoders.pkl", "wb") as f:
    pickle.dump(encoders, f)

```

In [32]:

encoders

Out[32]:

```

{'DVC_GRP': LabelEncoder(),
 'DVC_CLASS': LabelEncoder(),
 'DVC_SUPPORT': LabelEncoder(),
 'MOST_USED_4W_REGION': LabelEncoder(),
 'DTAC_RWRD_SEGMENT': LabelEncoder(),
 'PM_PMMTHDCOMMON': LabelEncoder()}

```

In [33]:

df\_fs.head()

Out[33]:

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_GRP	DVC_CLASS	DVC_SUPPORT	MOST_USED_4W_REGION	DTAC_RWRD_SEGMENT	Churn_F
0	1099.0	10519	467.0	1	1	5	0	1	
1	1099.0	10550	498.0	1	1	5	0	1	
2	1099.0	10580	528.0	1	1	5	0	1	
3	1099.0	10611	559.0	1	1	5	0	1	
4	999.0	10403	110.0	1	1	5	0	1	

5 rows × 22 columns



- Train test split

In [34]:

```

X = df_fs.drop(columns='Churn_Flag')
y = df_fs['Churn_Flag']

```

```
print(X.shape)
print(y.shape)
```

```
(178503, 21)
(178503,)
```

In [35]:

```
from sklearn.model_selection import train_test_split, cross_val_score

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print(X_train.shape)
print(y_train.shape)
```

```
(142802, 21)
(142802,)
```

- Synthetic Minority Oversampling Technique (SMOTE)

In [36]:

```
from imblearn.over_sampling import SMOTE

smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

print(y_train_smote.shape)
print(y_train_smote.value_counts())
```

```
(280790,)
0    140395
1    140395
Name: Churn_Flag, dtype: int64
```

- Model training

In [41]:

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

models = {
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "XGBoost": XGBClassifier(random_state=42)
}
```

```
In [42]: cv_scores = {}

for model_name, model in models.items():
    print(f"Training {model_name}")
    scores = cross_val_score(model, X_train_smote, y_train_smote, cv=5, scoring="accuracy")
    cv_scores[model_name] = scores
    print(f"{model_name} cross-validation accuracy: {np.mean(scores):.2f}")
    print("-"*70)
```

```
Training Decision Tree
Decision Tree cross-validation accuracy: 1.00
-----
```

```
Training Random Forest
Random Forest cross-validation accuracy: 1.00
-----
```

```
Training XGBoost
XGBoost cross-validation accuracy: 1.00
-----
```

```
In [43]: cv_scores
```

```
Out[43]: {'Decision Tree': array([0.99864668, 0.99960825, 0.99976851, 0.99985754, 0.99966167]),
          'Random Forest': array([0.99964386, 0.9997507 , 0.99982193, 0.99991097, 0.9997329 ]),
          'XGBoost': array([0.99966167, 0.99971509, 0.9997507 , 0.99982193, 0.99960825])}
```

```
In [44]: rf = RandomForestClassifier(random_state=42)
         rf.fit(X_train_smote, y_train_smote)
```

```
Out[44]: ▼      RandomForestClassifier
         RandomForestClassifier(random_state=42)
```

## 4. Model evaluation

```
In [45]: y_test_pred = rf.predict(X_test)

print("Accuracy Score:\n", accuracy_score(y_test, y_test_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_test_pred))
print("Classification Report:\n", classification_report(y_test, y_test_pred))
```

Accuracy Score:  
0.9994397916024761

Confusion Matrix:

```
[[35079    20]
 [     0   602]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	35099
1	0.97	1.00	0.98	602
accuracy			1.00	35701
macro avg	0.98	1.00	0.99	35701
weighted avg	1.00	1.00	1.00	35701

### To do:

1. Apply hyperparameter tuning to optimize model performance.
2. Experiment with different model architectures to identify the most effective one.
3. Use downsampling techniques to address class imbalance.
4. Implement strategies to reduce overfitting and improve generalization.
5. Utilize stratified K-Fold cross-validation for more reliable model evaluation.

## 5. Customer segmentation

- Categorized customers, who stay

In [51]:

```
df_st = df_fs[df_fs['Churn_Flag'] == 0]
df_st.head()
```

Out[51]:

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_GRP	DVC_CLASS	DVC_SUPPORT	MOST_USED_4W_REGION	DTAC_RWRD_SEGMENT	Churn_F
0	1099.0	10519	467.0	1	1	5	0	1	
1	1099.0	10550	498.0	1	1	5	0	1	
2	1099.0	10580	528.0	1	1	5	0	1	
3	1099.0	10611	559.0	1	1	5	0	1	
4	999.0	10403	110.0	1	1	5	0	1	

5 rows × 22 columns

```
In [52]: df_st.shape
```

```
Out[52]: (175494, 22)
```

```
In [56]: from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

X = df_st.drop(columns='Churn_Flag')

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

- Figure out an optimal point

```
In [58]: inertia = []
k_range = range(1, 11)
for k in k_range:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X_scaled)
    inertia.append(kmeans.inertia_)


# Plot the Elbow Method
plt.plot(k_range, inertia, marker='o')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Inertia')
plt.title('Elbow Method For Optimal k')
plt.show()
```

```
C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
```

```

value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)
C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)
C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)
C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)
C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)
C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)

```



```

In [61]: kmeans = KMeans(n_clusters=5, random_state=42)
df_st['Cluster'] = kmeans.fit_predict(X_scaled)

```

```

C:\Users\ADMIN\AppData\Roaming\Python\Python38\site-packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)
<ipython-input-61-081a115a96a4>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.

```



Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_st['Cluster'] = kmeans.fit_predict(X_scaled)
```

In [62]:

```
cluster_centers = pd.DataFrame(kmeans.cluster_centers_, columns=X.columns)
print("Cluster Centers (Centroids):")
print(cluster_centers)
```

Cluster Centers (Centroids):

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_GRP	DVC_CLASS	DVC_SUPPORT	\
0	-0.510445	-0.496990	-0.012716	-0.045965	0.163243	-0.061491	
1	1.168251	-0.285142	-0.165067	-0.055352	0.065311	0.263982	
2	1.668617	0.097881	-0.132899	-0.116135	0.054526	0.140319	
3	-0.130720	0.836246	0.051033	-0.087462	0.071069	0.110117	
4	-1.427695	-1.004254	1.289711	3.274363	-5.519996	-3.470822	

	MOST_USED_4W_REGION	DTAC_RWRD_SEGMENT	REVENUE	AVG3M_REVENUE	...	\
0	0.076695		0.962281	-0.522215	-0.531044	...
1	-0.150665		-0.427131	1.167230	1.183017	...
2	-0.007134		-0.211030	1.789458	1.820152	...
3	-0.016509		-0.995528	-0.130737	-0.128832	...
4	0.145196		1.013655	-1.382877	-1.438633	...

	VC_DOM_CNT	AVG3M_VC_DOM_MOU	AVG3M_VC_DOM_CNT	DATA_MB	AVG3M_DATA_MB	\
0	-0.270602	-0.267479	-0.271259	-0.286204	-0.286306	
1	0.206188	0.219908	0.209136	1.163129	1.177253	
2	3.491926	3.470099	3.492057	0.023961	0.021879	
3	-0.110546	-0.123284	-0.110368	-0.273698	-0.279691	
4	-0.669763	-0.614825	-0.688362	-0.760467	-0.792348	

	CIN_CALLCNT	VOC_ACTIVATEDAY	DATA_ACTIVATEDAY	PM_PMMTHDCOMMON	PCT_CALL_DROP
0	-0.069949	-0.281305	-0.253073	-0.055759	-0.089757
1	0.238471	0.438188	0.394719	-0.150501	0.280229
2	0.180587	1.180238	0.115629	0.244005	0.101554
3	-0.052597	0.088184	0.110715	0.081573	-0.040614
4	-0.293429	-1.972550	-0.593347	0.860256	-0.345078

[5 rows x 21 columns]

In [64]:

```
cluster_summary = df_st.groupby('Cluster').mean() # Average for each cluster
print("\nCluster Summary (Average values per cluster):")
print(cluster_summary)
```

Cluster Summary (Average values per cluster):

	CURR_MAIN_PKG_FEE	AOU_DAY	AOU_DVC	DVC_GRP	DVC_CLASS	\
--	-------------------	---------	---------	---------	-----------	---

Cluster	DVC_SUPPORT	MOST_USED_4W_REGION	DTAC_RWRD_SEGMENT	Churn_Flag	\
0	365.786732	2728.004712	620.082630	1.014341	1.022949
1	890.345556	3136.309154	543.762973	1.011574	1.005331
2	1046.892168	3871.231666	560.168852	0.993686	1.003400
3	484.368021	5290.843931	652.005568	1.002067	1.006388
4	79.205726	1753.207202	1272.420897	2.000000	0.000000

Cluster	REVENUE	...	VC_DOM_CNT	AVG3M_VC_DOM_MOU	AVG3M_VC_DOM_CNT	\
0	365.562444	...	48.240425	110.423989	49.272431	
1	941.231515	...	103.491400	258.974407	103.826579	
2	1153.402936	...	484.199126	1249.517997	476.497949	
3	498.843061	...	66.749398	154.305238	67.504372	
4	72.341502	...	2.030401	4.657713	1.960449	

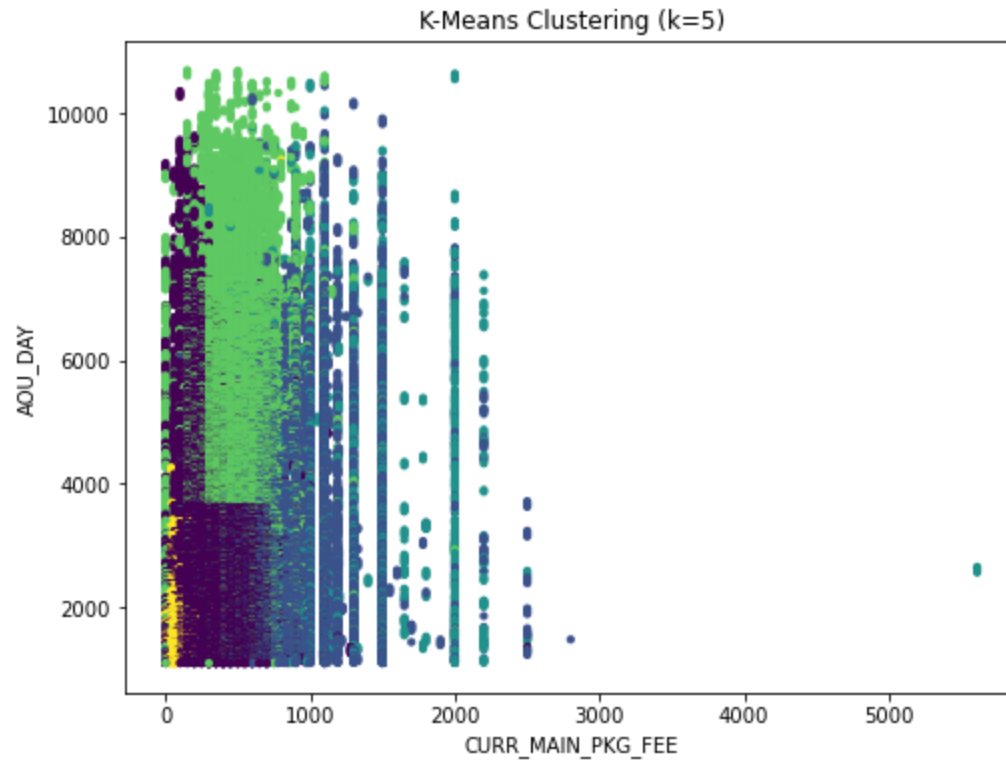
Cluster	DATA_MB	AVG3M_DATA_MB	CIN_CALLCNT	VOC_ACTIVATEDAY	\
0	12011.480431	12335.799479	0.905576	16.164919	
1	47843.794290	46911.936373	2.094599	22.850592	
2	19693.985545	19626.491304	1.874373	29.738546	
3	12318.506255	12489.781942	0.972725	19.594524	
4	275.739687	371.205955	0.043388	0.457202	

Cluster	DATA_ACTIVATEDAY	PM_PMMTHDCOMMON	PCT_CALL_DROP
0	24.455587	1.878214	0.416897
1	30.130500	1.612829	1.019842
2	27.694026	2.718472	0.729498
3	27.640706	2.263070	0.496850
4	21.473731	4.441558	0.000000

[5 rows x 22 columns]

```
In [68]: plt.figure(figsize=(8,6))
plt.scatter(df_st.iloc[:, 0], df_st.iloc[:, 1], c=df_st['Cluster'], cmap='viridis', s=10)
plt.title('K-Means Clustering (k=5)')
plt.xlabel(X.columns[0])
plt.ylabel(X.columns[1])
plt.show()
```

**To do:**

1. Tailor benefits and offers for each customer segment based on their specific needs and behavior.
2. Monitor and evaluate the effectiveness of each offer or strategy using measurable KPIs.
3. Identify which offers drive positive results, and proactively present those to customers who are at high risk of churning.